




Every learner degrades to “no opinion” at zero data rather than to a confident guess: an unfitted head returns nothing and the judge ranks instead, which is what makes the no-preference ablation exact. Learned state is persisted per branch, so taste explored on a what-if branch cannot leak back to the main line.